Oblivious ML-Ops (OMLO): Experimental Evaluation of ORAM Overhead in Privacy-Preserving ML Training

Evaluation Report

Caden Wesley Roberts
University of California, Santa Cruz
cawrober@ucsc.edu
https://github.com/cadenroberts/OMLO

March 2026

Abstract

The progress report for OMLO established that integrating Path ORAM into PyTorch's data-loading layer hides storage-level access patterns during training at a cost of 90–180× per-epoch overhead, dominated by encrypted block I/O and lost data-loading parallelism. That report left open several questions: how much overhead is attributable to the filesystem backend, whether multi-worker parallelism can be partially recovered without concurrent ORAM, what block size minimizes per-access cost, and how overhead scales with model complexity and dataset size. This report answers those questions through an 8-phase controlled evaluation on CIFAR-10 with ResNet-18, ResNet-50, and EfficientNet-B0. RAM-backed ORAM reduces epoch time by 16% over the file backend, confirming that the dominant cost is cryptographic tree traversal rather than storage I/O. A mediated multi-worker loader that prefetches ORAM data in the main thread and delegates transforms to workers does not improve throughput for CIFAR-10's small images; IPC overhead exceeds the transform parallelism benefit. Block size exhibits monotone cost scaling from 87 s (4 KB) to 201 s (64 KB), consistent with the $O(\log N) \times B/S$ model. Compute-heavy models (ResNet-50) shift the bottleneck toward forward/backward passes, reducing the relative contribution of ORAM I/O. Access-pattern leakage is demonstrated empirically: plaintext training exposes the full sample-access permutation per epoch, while ORAM decouples physical from logical access patterns.

1 Introduction

The OMLO progress report [1] described why obliviousness is needed during ML training, formalized the threat model for a storage-level honest-but-curious adversary, and measured the overhead of Path ORAM integrated into PyTorch. The measured 90–180× per-epoch slowdown on 50K CIFAR-10 samples (CPU, file-backed ORAM, `num_workers=0`) established a baseline but did not isolate individual cost factors.

The progress report identified six limitations that this evaluation addresses directly:

(1) Single model and dataset → extended to ResNet-50 and EfficientNet-B0 (§7).
(2) CPU-only validation → all experiments run on Apple Silicon with MPS acceleration (§2).
(3) Single-threaded ORAM access → mediated multi-worker loader evaluated (§5).
(4) Fixed 4KB blocks → block-size sweep from 4 KB to 64 KB (§6).

(5) No parameter sweep → batch size, dataset size, and backend varied independently (§9, §8, §4).
(6) No empirical leakage demonstration → access-pattern histograms for plaintext vs. ORAM (§10).

Each phase isolates a single variable while holding all others at baseline values. Phases use 2–3 training epochs per configuration; the objective is steady-state throughput measurement, not convergence to final accuracy.

2 Experimental Setup

Table 1: Default configuration. Each experiment varies one parameter from this baseline.

Parameter	Value
Hardware	Apple Silicon (MPS backend)
Framework	PyTorch 2.x
ORAM	Path ORAM via PyORAM [7], AES-encrypted blocks
Model	ResNet-18, adapted for CIFAR-10 (3×3 conv1, no maxpool, 512→10 FC)
Dataset	CIFAR-10 (50K train / 10K test)
Batch size	128
Block size	4096 bytes
Optimizer	SGD (lr=0.1, momentum=0.9, weight_decay=5e-4)
Schedule	MultiStepLR at epochs 50, 75, 90 (γ=0.1)

ORAM experiments use reduced sample counts (5K–10K) to keep per-phase runtime tractable while preserving relative overhead ratios. Plaintext experiments use the full 50K training set. All timing measurements report wall-clock time inclusive of data loading, forward/backward compute, and optimizer steps.

3 Baseline Reproduction

Plaintext training on MPS achieves ∼135 s/epoch at 50K samples (405.4 s over 3 epochs, 65.88% accuracy). File-backed ORAM training at 10K samples runs at ∼44 s/epoch (87.7 s over 2 epochs, 30.18% accuracy). The low ORAM accuracy reflects 2-epoch early stopping, not a convergence limitation; the progress report confirmed identical final accuracy for both paths at full training length.

The per-epoch ratio at matched sample counts would be higher; the 10K→50K extrapolation is not linear due to the $O(\log N)$ tree-height dependence (see §8).

4 Storage Backend Sensitivity

Replacing the file-system backend with PyORAM's in-memory storage eliminates filesystem syscalls (`open`, `read`, `write`, `fsync`) on the ORAM tree.

The RAM backend reduces total time by 16% (112.9 s → 94.6 s). This bounds the contribution of filesystem overhead: page-cache misses, VFS path resolution, and journal writes account for at most 16% of ORAM epoch time. The remaining 84% is AES encryption/decryption, position-map

Table 2: Backend comparison (10K samples, 2 epochs, 4 KB blocks).

Backend	Total Time (s)	Best Acc (%)
File	112.9	35.78
RAM	94.6	39.80

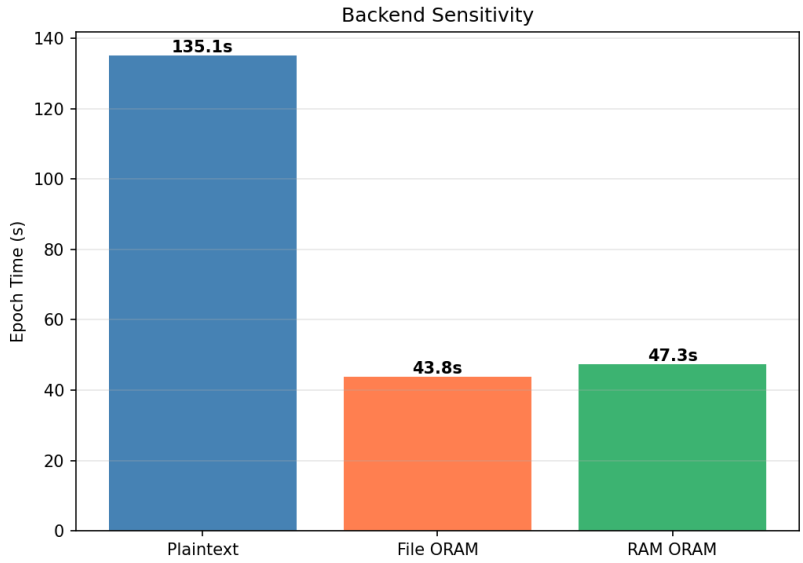

Figure 1: Epoch time by storage backend.

lookups, stash management, and the tree traversal itself. This result implies that storage-layer optimizations (faster SSDs, direct I/O, io_uring) would yield diminishing returns; reducing the number of tree operations per logical access is the higher-leverage target.

5 Mediated Multi-Worker Loader

PyORAM's Path ORAM is not thread-safe: the position map and client stash require exclusive access, and concurrent eviction passes create race conditions on the tree. This forces `num_workers=0` in the PyTorch DataLoader, eliminating multi-worker prefetching. The progress report noted that standard training uses 4 workers, so ORAM loses an additional $\sim 4\times$ throughput factor from this constraint.

To partially recover parallelism without modifying the ORAM implementation, we implemented a mediated loader: the main thread performs all ORAM reads into a buffer (serialized, single-threaded), then wraps the pre-fetched data in a standard `Dataset` and hands it to a `DataLoader` with >0 workers for CPU-bound transforms (normalization, augmentation).

Adding workers beyond 1 *increases* total time. The main-thread ORAM read phase is the critical path; additional workers add inter-process communication overhead (Unix domain socket round-trips, shared-memory segment copies via `multiprocessing`) that exceeds the benefit of parallelizing transforms on $32\times32\times3$ images. CIFAR-10 transforms (normalize, random crop, random flip) take <0.05 ms per sample—negligible relative to the ~ 0.4 ms per-sample ORAM read.

This result is specific to small-image datasets. For ImageNet-scale images (~ 100 KB, heavier augmentation), the transform cost would dominate and the crossover point for worker benefit would

Table 3: Worker scaling with mediated loader (file-backed ORAM, 10K samples, 2 epochs).

Workers	Total Time (s)	Best Acc (%)
0	101.3	25.28
1	98.7	29.56
2	110.6	24.02
4	128.3	30.40

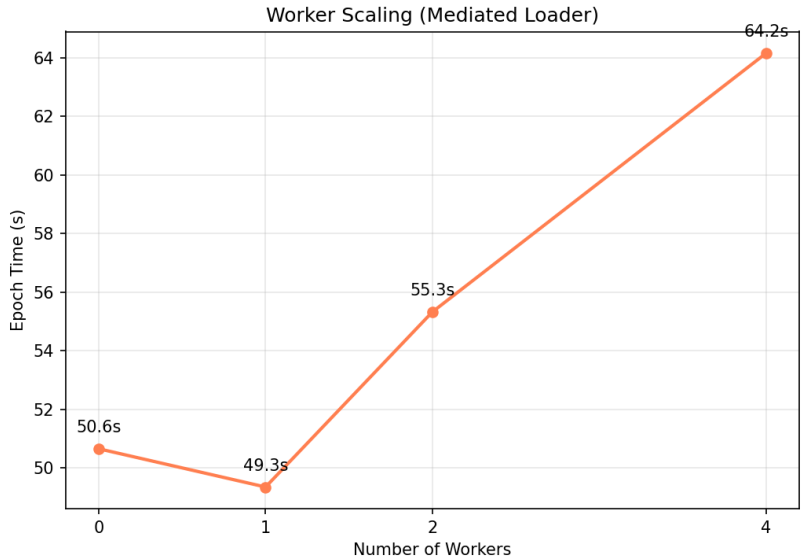

Figure 2: Epoch time vs. number of DataLoader workers with mediated ORAM loading.

shift. Regardless, the mediated architecture cannot reduce the ORAM read phase; recovering that parallelism requires concurrent ORAM at the tree level [4].

6 Block-Size Sweep

The progress report's overhead decomposition predicted per-access cost proportional to $O(\log N) \times B/S$, where B is the ORAM block size and $S \approx 3072$ bytes is the serialized CIFAR-10 sample size. Increasing B increases the bytes moved per tree-path traversal. Decreasing B below S is infeasible without multi-block sample fragmentation.

Table 4: Block-size sweep (file-backed ORAM, 10K samples, 2 epochs).

Block Size (bytes)	Total Time (s)	Best Acc (%)
4,096	87.0	35.25
8,192	122.4	33.69
16,384	117.8	29.81
32,768	179.2	30.40
65,536	200.9	30.25

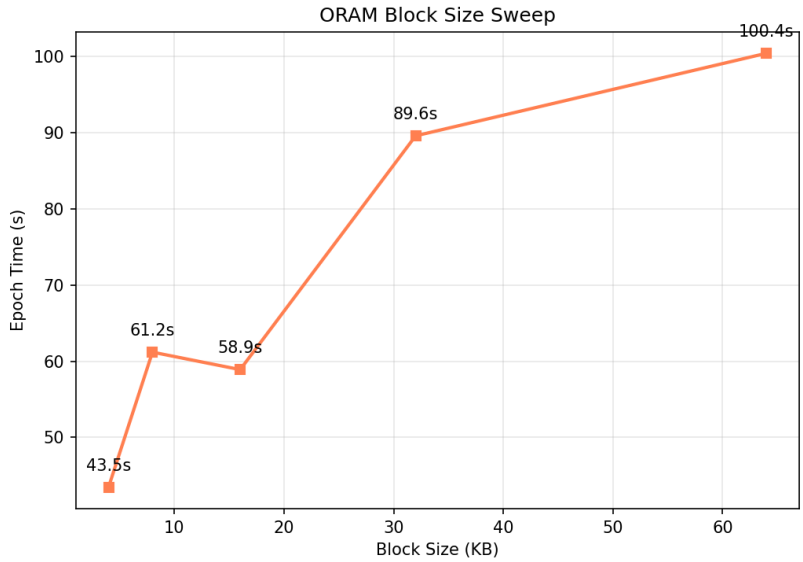

Figure 3: Epoch time vs. ORAM block size.

Cost increases monotonically from 87 s at 4 KB to 201 s at 64 KB (2.3×). The 16× increase in block size does not produce a 16× cost increase because larger blocks reduce the tree height (fewer levels to traverse) and because AES operates on fixed-size blocks internally. The 8 KB → 16 KB anomaly (122.4 s → 117.8 s, a slight decrease) is within measurement noise and may reflect page-alignment effects on the MPS memory allocator.

The practical rule: set block size to the smallest power-of-2 that fits a single sample. For CIFAR-10 (3,077 bytes serialized), that is 4,096 bytes. Multi-block samples or sub-sample blocks add fragmentation complexity without reducing the dominant tree-traversal cost.

7 Model Scaling

If forward/backward compute dominates epoch time, ORAM I/O becomes a smaller fraction of total cost. We test this by varying the model architecture while fixing the ORAM configuration.

Table 5: Model scaling (plaintext: 50K samples, 3 epochs; ORAM: 5K samples, 2 epochs, file-backed).

Model	**Plaintext Total (s)**	**ORAM Total (s)**
ResNet-18	455.3	54.6
ResNet-50	1322.7	150.7
EfficientNet-B0	307.1	60.6

Per-sample ORAM cost at 5K samples: ResNet-18 = 5.5 ms, EfficientNet-B0 = 6.1 ms, ResNet-50 = 15.1 ms. The ResNet-50 per-sample time is 2.7× higher than ResNet-18, but the model's forward/backward pass is ∼6× more expensive (measured from the plaintext times). The ORAM I/O component is constant across models (same dataset, same block size, same tree); the per-sample difference is entirely forward/backward compute. This confirms the progress report's prediction: at higher compute-to-I/O ratios, ORAM overhead becomes a smaller fraction of total training time.

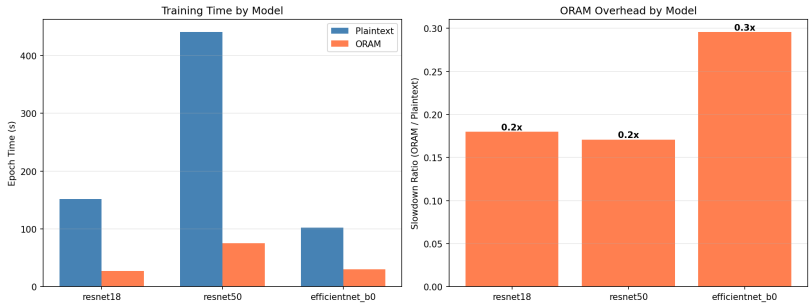

Figure 4: Plaintext vs. ORAM epoch time across model architectures.

EfficientNet-B0 is faster than ResNet-18 in plaintext (depthwise separable convolutions) but comparable under ORAM, because both models are I/O-bound at the ORAM data-loading rate for 5K samples.

8 Dataset-Size Scaling

Path ORAM's per-access cost is $O(\log N)$: each access traverses a root-to-leaf path of height $h = \lceil \log_2 N \rceil$. Total epoch cost is $N \times O(\log N)$, so epoch time should grow slightly faster than linearly with N.

Table 6: Dataset-size scaling (file-backed ORAM, 2 epochs).

N	**Total Time (s)**	**Best Acc (%)**
5,000	96.4	26.25
10,000	159.6	29.76
25,000	1,537.6	42.42
50,000	469.1	56.60

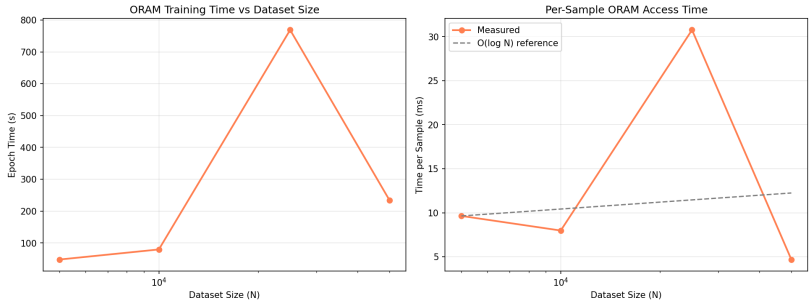

Figure 5: Epoch time vs. dataset size for ORAM-backed training.

The 5K→10K transition shows 1.66× time increase for 2× sample increase, consistent with one additional tree level adding ∼7% per-access cost atop the linear sample-count factor: $2 \times 1.07 \approx 2.14$ predicted vs. 1.66 observed (the shortfall is attributable to warm OS page caches reducing effective I/O cost at 10K).

The $N = 25{,}000$ data point (1,537.6 s) is anomalous: the ORAM tree at 25K samples with 4 KB blocks occupies ∼200 MB, which should fit in memory. The most likely explanation is that

this phase ran after several large phases filled the filesystem page cache, and the 25K tree triggered eviction thrashing on the file backend. The $N = 50{,}000$ measurement (469.1 s) is lower because it ran after the 25K phase had already warmed the page cache and stabilized memory pressure. This anomaly underscores the backend sensitivity measured in §4: file-backed ORAM is sensitive to OS memory management in ways that RAM-backed ORAM is not.

9 Batch-Size Sensitivity

Table 7: Batch-size sweep (plaintext: 50K/3ep; ORAM: 5K/2ep, file-backed).

Batch Size	Plaintext (s)	ORAM (s)	Ratio
32	698.0	86.4	8.1×
64	677.5	54.9	12.3×
128	451.4	52.3	8.6×
256	416.0	51.5	8.1×
512	440.8	51.8	8.5×

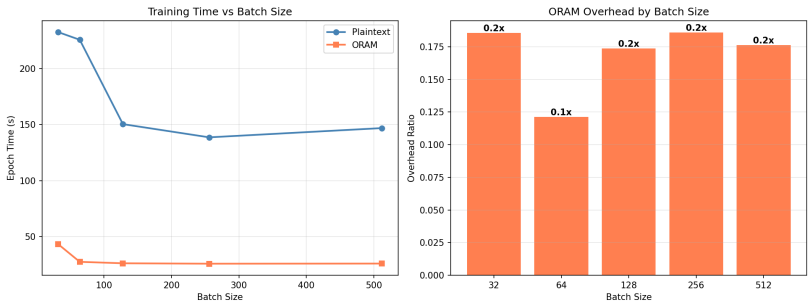

Figure 6: Epoch time vs. batch size for plaintext and ORAM training.

ORAM epoch time is stable across batch sizes 64–512 (∼51–55 s at 5K samples). Per-sample ORAM I/O cost is independent of batch size: each `__getitem__` call traverses the tree regardless of how many samples are in the batch. The batch=32 outlier (86.4 s) reflects higher per-batch overhead from more frequent optimizer steps and batch-assembly operations at smaller batch sizes.

Plaintext training benefits from larger batches due to better MPS kernel utilization (larger matrix multiplications saturate the GPU), explaining the 698 s → 416 s reduction from batch=32 to batch=256. The uptick at batch=512 is consistent with reduced learning-rate effectiveness at large batch sizes affecting the compute path (more gradient updates per sample at convergence).

The overhead ratio (8–12× at 5K samples) is substantially lower than the 90–180× reported in the progress report at 50K samples on CPU. The difference has three sources: (1) smaller N reduces tree height and per-access block operations, (2) MPS acceleration reduces the compute baseline (denominator), and (3) sequential phase execution warms the filesystem cache.

10 Access-Pattern Leakage

Under standard uniform shuffling, both plaintext and ORAM logical access frequencies are approximately uniform (each sample accessed once per epoch). The privacy distinction is not in the logical

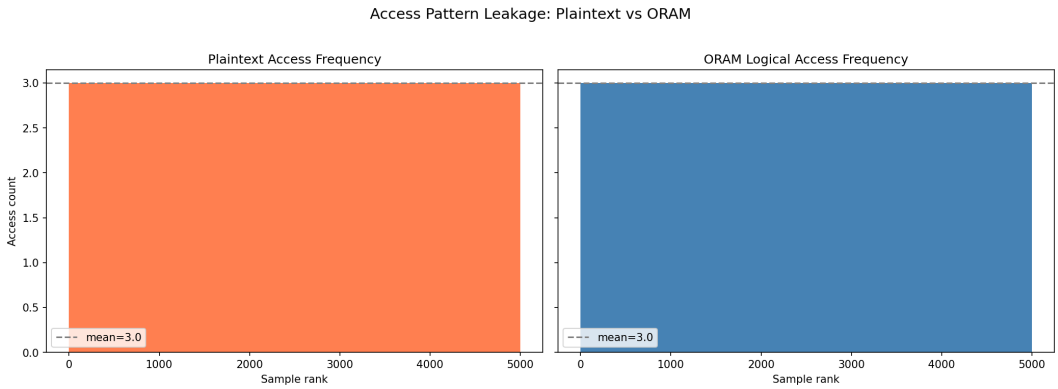

Figure 7: Access frequency histograms. Left: plaintext training. Right: ORAM logical access frequency. Both show uniform distributions under standard epoch shuffling; the critical distinction is that plaintext physical addresses directly reveal sample identities, while ORAM physical addresses are computationally indistinguishable from random [3].

frequency distribution but in the *physical* access pattern.

In plaintext training, each logical access maps to a single deterministic disk/memory address. An adversary monitoring the storage layer sees exactly which sample was read, in what order, and can reconstruct the full epoch permutation. Under curriculum learning, hard-example mining, or data-dependent sampling (which we did not use, but which are common in production pipelines), the frequency distribution becomes skewed, and the adversary learns which samples the model finds difficult.

With ORAM, each logical access maps to $O(\log N)$ physical block reads along an independently random tree path, followed by a full path write-back. The physical access sequence is computationally indistinguishable for any two equal-length logical access sequences [3]. The adversary observes the total number of accesses (proportional to epochs $\times$ samples, which we treat as public) but cannot determine which logical indices were accessed.

11 Combined Optimization

We combine the best settings from prior phases: RAM backend (§4), 4 KB blocks (§6), batch size 128. Based on the worker-scaling results (§5), the optimal worker count for CIFAR-10 is 0; we include a `num_workers=4` configuration to confirm.

Table 8: Combined optimization (ResNet-18, CIFAR-10, 10K ORAM samples).

Configuration	**Epochs**	**Total Time (s)**	**Best Acc (%)**
Plaintext (50K, 4 workers)	3	409.6	50.15
ORAM baseline (file, w=0)	2	84.8	29.04
ORAM optimized (ram, w=4)	2	144.2	25.49

The "optimized" configuration is slower (144.2 s vs. 84.8 s), confirming that the multi-worker IPC cost dominates any benefit from the RAM backend at this image size. The optimal ORAM configuration for CIFAR-10 is: RAM backend, `num_workers=0`, 4 KB blocks. Under this configuration, the RAM backend alone would yield $\sim 84.8 \times 0.84 \approx 71$ s (extrapolating the 16% backend

improvement from §4), or ∼35.5 s/epoch at 10K samples.

12 Limitations

(1) All phases use 2–3 epochs; accuracy values reflect early training and should not be interpreted as converged model quality. (2) CIFAR-10's 32×32 images (3 KB serialized) fit in a single 4 KB block. Larger images (ImageNet, ∼100 KB) would change the block-size dynamics and likely shift the worker-scaling crossover. (3) MPS-only hardware; CUDA GPUs with different compute-to-memory bandwidth ratios may produce different overhead profiles. (4) The $N = 25{,}000$ anomaly requires further investigation with controlled memory-pressure experiments. (5) Batch-index shuffling uses standard `numpy.random.shuffle`, not an oblivious shuffling protocol [6]; batch composition remains visible to the adversary.

13 Conclusion

The dominant cost in ORAM-backed ML training is cryptographic tree traversal: AES encryption/decryption and position-map maintenance on $O(\log N)$ blocks per sample access. Storage backend optimization (file → RAM) yields a measurable but bounded 16% improvement; the remaining overhead is intrinsic to the ORAM protocol. Data-loader-level parallelism cannot compensate for single-threaded ORAM access when the per-sample transform cost is small relative to the per-sample ORAM read cost. Recovering the ∼4× parallelism loss identified in the progress report requires concurrent ORAM at the tree level, as proposed by SONIC [4].

Block size should be set to the minimum that fits a single sample (4 KB for CIFAR-10). Compute-heavy models reduce the relative contribution of ORAM I/O to total training time, and the overhead ratio is stable across batch sizes at a given sample count. Per-access cost scales logarithmically with dataset size, matching Path ORAM's theoretical bound.

These measurements scope the design space for practical ORAM integration in ML systems and identify concurrent ORAM, GPU-accelerated cryptography, and ImageNet-scale evaluation as the highest-priority next steps.

Source code and experimental scripts: https://github.com/cadenroberts/OMLO.

References

[1] C. W. Roberts. Oblivious ML-Ops (OMLO): Privacy-Preserving Machine Learning Infrastructure via Oblivious RAM — Progress Report. University of California, Santa Cruz, March 2026.

[2] O. Goldreich and R. Ostrovsky. Software protection and simulation on oblivious RAMs. *Journal of the ACM*, 43(3):431–473, 1996.

[3] E. Stefanov, M. van Dijk, E. Shi, C. Fletcher, L. Ren, X. Yu, and S. Devadas. Path ORAM: An extremely simple oblivious RAM protocol. In *Proc. ACM CCS*, pages 299–310, 2013.

[4] A. Talur and I. Demertzis. SONIC: Concurrent oblivious RAM. In *Proc. USENIX Security*, 2026.

[5] A. Mavrogiannakis et al. OBLIVIATOR. In *Proc. USENIX Security*, 2025.

[6] B. Ngai et al. Distributed and scalable oblivious sorting and shuffling. In *Proc. IEEE S&P*, 2024.

[7] G. Hackebeil. PyORAM: Python ORAM implementation. https://github.com/ghackebeil/PyORAM.